

▼ Problem Statement

- A dataset collected in a cosmetics shop showing details of DEFINITION customers and whether or not they responded to a special offer to buy a new lip-stick is shown in table below.
- Use this dataset to build a decision tree, with Buys as the target variable, to help in buying lip-sticks in the future.
- Find the root node of decision tree.
- According to the decision tree you have made from previous training data set, what is the decision for the test data:
 - [Age < 21, Income = Low, Gender = Female, Marital Status = Married]?

ID	Age	Income	Gender	Marital Status	Buys
1	<21	High	M	Single	N
2	<21	High	M	Married	N
3	21-35	High	M	Single	Y
4	>35	Medium	M	Single	Y
5	>35	Low	F	Single	Y
6	>35	Low	F	Married	N
7	21-35	Low	F	Married	Y
8	<21	Medium	M	Single	N
9	<21	Low	F	Married	Y
10	>35	Medium	F	Single	Y
11	<21	Medium	F	Married	Y
12	21-35	Medium	M	Married	Y
13	21-35	High	F	Single	Y
14	>35	Medium	M	Married	N

Gini Impurity Formula

$$Gini = 1 - \sum_{i=1}^C (p_i)^2$$

Information Gain Formula

- $p = \text{size_of_true_rows} / (\text{size_of_true_rows} + \text{size_of_false_rows})$
- $\text{current_uncertainty} = \text{gini impurity of input rows}$

- Compute the weighted average of the gini impurity of the true branch and the false branch, and subtract it from the current uncertainty

$$\text{information_gain} = \text{current_uncertainty} - p * \text{gini}(\text{true_rows}) - (1 - p) * \text{gini}(\text{false_rows})$$

Decision Tree

- Full binary tree - each node contains either 2 or 0 child nodes
 - Internal Nodes - Question
 - Leaf Nodes - Final Prediction
 - Edges - Answer to the Question
-

CART Algorithm

Algorithm: build_tree(rows)

- rows is the rows of the training_data passed to build_tree()
- find_best_split(rows) is used to determine the best question to ask
- Leaf is a class which returns final prediction
- partition(rows, question) is used to split the input rows based on the question asked
- Decision_Node is a class which contains the question, the true_branch and the false_branch

Steps

1. information_gain, question = find_best_split(rows)
 2. if information_gain := 0 then
 1. return Leaf(rows)
 3. true_rows, false_rows = partition(rows, question)
 4. true_branch = build_tree(true_rows)
 5. false_branch = build_tree(false_rows)
 6. return Decision_Node(question, true_branch, false_branch)
-

Algorithm: find_best_split(rows)

- best_question is used to keep a track of the question with the highest information gain.
- best_gain is the information gain of the best_question.

- `current_uncertainty` is the gini impurity of the input rows.
- `Question` is a class which contains the `input_feature` and one of its unique values.
- `partition` is the function which splits the input rows into `true_rows` and `false_rows` based on the question asked.

Steps

1. `best_question = None`
2. `best_gain = 0`
3. for each `input_feature` in `input_features` do
 1. `unique_values = Get a set of unique values in the feature`
 2. for each `unique_value` in `unique_values`
 - 3.2.1. `question = Question(input_feature, unique_value)`
 - 3.2.2. `true_rows, false_rows := partition(rows, question)`
 - 3.2.3. If question does not split the dataset then skip it
 - 3.2.4. `information_gain := gain(true_rows, false_rows, current_uncertainty)`
 - 3.2.5. if `information_gain > best_gain` then:
 - 3.2.5.1 `best_gain := information_gain`
 - 3.2.5.2 `best_question := question`
4. return `best_question, best_gain`

▼ Source Code

```

1 # Friend Functions
2
3 # Find the unique values for a column in a dataset.
4 def unique_vals(rows, col):
5     return set([row[col] for row in rows])
6
7
8 # Counts the number of each type of example in a dataset.
9 def class_counts(rows):
10     counts = {} # a dictionary of label -> count.
11     for row in rows:
12         # in our dataset format, the label is always the last column
13         label = row[-1]
14         if label not in counts:
15             counts[label] = 0
16         counts[label] += 1

```

```

16         counts[label] += 1
17     return counts
18
19 #####
20
21 # A Question is used to partition a dataset.
22 # This class just records a 'column number' and a 'column value'.
23 # The 'match' method is used to compare the feature value
24 # in an example to the feature value stored in the question.
25 class Question:
26
27
28     def __init__(self, column, value):
29         self.column = column
30         self.value = value
31
32
33     # Compare the feature value in an example to the
34     # feature value in this question.
35     def match(self, example):
36         val = example[self.column]
37         return val == self.value
38
39
40     def __repr__(self):
41         # This is just a helper method to print
42         # the question in a readable format.
43         return "is {column} == {value}".format( column=header[self.column], value=
44
45 # A Leaf node classifies data.
46 # This holds a dictionary of class (e.g., "Apple") -> number of times
47 # it appears in the rows from the training data that reach this leaf.
48 class Leaf:
49
50     def __init__(self, rows):
51         self.predictions = class_counts(rows)
52
53 # A Decision Node asks a question.
54 # This holds a reference to the question, and to the two child nodes.
55 class Decision_Node:
56
57     def __init__(self,
58                 question,
59                 true_branch,
60                 false_branch):
61         self.question = question
62         self.true_branch = true_branch
63         self.false_branch = false_branch
64
65
66 #####
67

```

```

68 class Decision_Tree:
69
70
71     # Partitions a dataset.
72     # For each row in the dataset, check if it matches the question.
73     # If so, add it to 'true rows', otherwise, add it to 'false rows'.
74     def partition(self, rows, question):
75         true_rows, false_rows = [], []
76         for row in rows:
77             if question.match(row):
78                 true_rows.append(row)
79             else:
80                 false_rows.append(row)
81         return true_rows, false_rows
82
83
84     # Calculate the Gini Impurity for a list of rows.
85     def gini(self, rows):
86         counts = class_counts(rows)
87         impurity = 1
88         for lbl in counts:
89             prob_of_lbl = counts[lbl] / float(len(rows))
90             impurity -= prob_of_lbl**2
91         return impurity
92
93
94     # Information Gain. - The uncertainty of the starting node, minus the weighted
95     def info_gain(self, left, right, current_uncertainty):
96         p = float(len(left)) / (len(left) + len(right))
97         return current_uncertainty - p * self.gini(left) - (1 - p) * self.gini(right)
98
99
100     # Find the best question to ask by iterating over every feature / value
101     # and calculating the information gain.
102     def find_best_split(self, rows):
103         best_gain = 0 # keep track of the best information gain
104         best_question = None # keep track of the feature / value that produced it
105         current_uncertainty = self.gini(rows)
106         n_features = len(rows[0]) - 1 # number of columns
107
108         # For each feature
109         for col in range(n_features):
110
111             # unique values in the column
112             values = set([row[col] for row in rows])
113
114             # for each value
115             for val in values:
116
117                 # Ask Question
118                 question = Question(col, val)
119

```

```

119
120         # Split the dataset
121         true_rows, false_rows = self.partition(rows, question)
122
123         # Skip this split if it doesn't divide the dataset.
124         if len(true_rows) == 0 or len(false_rows) == 0:
125             continue
126
127         # Calculate the information gain from this split
128         gain = self.info_gain(true_rows, false_rows, current_uncertainty)
129
130         # Record the best gain and the best question
131         if gain >= best_gain:
132             best_gain, best_question = gain, question
133
134     # Return Question with the highest information gain
135     return best_gain, best_question
136
137
138 # Builds the tree.
139 # Rules of recursion:
140 # 1) Assume that it works.
141 # 2) Start by checking for the base case (no further information gain).
142 # 3) Prepare for giant stack traces.
143 def build_tree(self, rows):
144
145     # Try partitioning the dataset on each of the unique attribute,
146     # calculate the information gain,
147     # and return the question that produces the highest gain.
148     gain, question = self.find_best_split(rows)
149
150     # Base case: no further info gain
151     # Since we can ask no further questions,
152     # we'll return a leaf.
153     if gain == 0:
154         return Leaf(rows)
155
156     # If we reach here, we have found a useful feature / value
157     # to partition on.
158     true_rows, false_rows = self.partition(rows, question)
159
160     # Recursively build the true branch.
161     true_branch = self.build_tree(true_rows)
162
163     # Recursively build the false branch.
164     false_branch = self.build_tree(false_rows)
165
166     # Return a Question node.
167     # This records the best feature / value to ask at this point,
168     # as well as the branches to follow
169     # depending on the answer.
170     return Decision_Node(question, true_branch, false_branch)

```

```

171
172
173 def print_tree(self,node, spacing=""):
174
175     # Base case: we've reached a leaf
176     if isinstance(node, Leaf):
177         print (spacing + "Predict", node.predictions)
178         return
179
180     # Print the question at this node
181     print (spacing + str(node.question))
182
183     # Call this function recursively on the true branch
184     print (spacing + '--> True:')
185     self.print_tree(node.true_branch, spacing + " ")
186
187     # Call this function recursively on the false branch
188     print (spacing + '--> False:')
189     self.print_tree(node.false_branch, spacing + " ")
190
191
192 # Rules of recursion:
193 # 1) Assume that it works.
194 # 2) Start by checking for the base case (no further information gain).
195 # 3) Prepare for giant stack traces.
196 def classify(self,row, node):
197
198     # Base case: we've reached a leaf
199     if isinstance(node, Leaf):
200         return node.predictions
201
202     # Decide whether to follow the true-branch or the false-branch.
203     # Compare the feature / value stored in the node,
204     # to the example we're considering.
205     if node.question.match(row):
206         return self.classify(row, node.true_branch)
207     else:
208         return self.classify(row, node.false_branch)
209
210
211 def print_leaf(self,counts):
212     total = sum(counts.values()) * 1.0
213     probs = {}
214     for lbl in counts.keys():
215         probs[lbl] = str(int(counts[lbl] / total * 100)) + "%"
216     return probs
217
218 def set_root(self,node):
219     self.root = node
220
221 def get_root(self):
222     return self.root

```

```

222         return self.root
223
224 #####
225 class PCAG:
226
227
228     def __init__(self, training_data, testing_data, header):
229         training_data = self.data_validation(training_data, header)
230         DTC = self.training(training_data)
231         self.prediction(DTC, testing_data)
232
233
234     def data_validation(self, training_data, header):
235
236         print('\nData Validation\n=====')
237
238         print('# Raw Train Data = {length}'.format(length=len(training_data)))
239         training_data_cleaned = []
240         for record in training_data:
241             age, income, gender, marital_status, buys = record
242             if age not in ['<21', '21-35', '>35']:
243                 pass
244             if income not in ['Low', 'Medium', 'High']:
245                 pass
246             if gender not in ['M', 'F']:
247                 pass
248             if marital_status not in ['Single', 'Married']:
249                 pass
250             if buys not in ['N', 'Y']:
251                 pass
252             training_data_cleaned.append(
253                 [age, income, gender, marital_status, buys]
254             )
255         print('# Cleaned Train Data = {length}'.format(length=len(training_data_cleaned)))
256
257         return training_data_cleaned
258
259
260     def training(self, training_data):
261         print('\nTraining\n=====')
262         DTC = Decision_Tree()
263         decision_tree = DTC.build_tree(training_data)
264         DTC.set_root(decision_tree)
265         print('Final Tree')
266         DTC.print_tree(DTC.get_root())
267         return DTC
268
269
270     def prediction(self, DTC, testing_data):
271         print('\nPrediction\n=====')
272         print('Prediction')
273         for row in testing_data:

```



```

274     print("Record: {record}".format(record=row))
275     print("Predicted: {predicted}".format(predicted=DTC.print_leaf(DTC.classify(
276     pass
277
278 #####
279
280 training_data = [
281     ['<21', 'High', 'M', 'Single', 'N'],
282     ['<21', 'High', 'M', 'Married', 'N'],
283     ['21-35', 'High', 'M', 'Single', 'Y'],
284     ['>35', 'Medium', 'M', 'Single', 'Y'],
285     ['>35', 'Low', 'F', 'Single', 'Y'],
286     ['>35', 'Low', 'F', 'Married', 'N'],
287     ['21-35', 'Low', 'F', 'Married', 'Y'],
288     ['<21', 'Medium', 'M', 'Single', 'N'],
289     ['<21', 'Low', 'F', 'Married', 'Y'],
290     ['>35', 'Medium', 'F', 'Single', 'Y'],
291     ['<21', 'Medium', 'F', 'Married', 'Y'],
292     ['21-35', 'Medium', 'M', 'Married', 'Y'],
293     ['21-35', 'High', 'F', 'Single', 'Y'],
294     ['>35', 'Medium', 'M', 'Married', 'N']
295 ]
296 testing_data = [
297     ['<21', 'Low', 'F', 'Married']
298 ]
299 header = ["Age", "Income", "Gender", "Marital Status", "Buys"]
300 pcag = PCAG(training_data, testing_data, header)

```

Data Validation

=====

Raw Train Data = 14

Cleaned Train Data = 14

Training

=====

Final Tree

is Age == 21-35

--> True:

Predict {'Y': 4}

--> False:

is Gender == F

--> True:

is Marital Status == Married

--> True:

is Age == >35

--> True:

Predict {'N': 1}

--> False:

Predict {'Y': 2}

--> False:

Predict {'Y': 2}

--> False:

is Age == >35

```
--> True:
  is Marital Status == Married
--> True:
  Predict {'N': 1}
--> False:
  Predict {'Y': 1}
--> False:
  Predict {'N': 3}
```

Prediction

=====

Prediction

Record: ['<21', 'Low', 'F', 'Married']

Predicted: {'Y': '100%'}

